

Problem 2 - OpenML, Algorithmic Performance Scaling (25 points)

This notebook explores classification tasks using datasets from OpenML, comparing Random Forest and Gradient Boosting classifiers.

Setup: Import Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import time
import warnings
from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier, GradientBoostingC
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import LabelEncoder

warnings.filterwarnings('ignore')

# Set random seed for reproducibility
np.random.seed(42)

print("Libraries imported successfully.")
```

Libraries imported successfully.

Part 1: Load and Summarize Datasets [5 points]

Task: Select 2 datasets from OpenML with different number of output classes and summarize their attributes.

```
In [31]: # Dataset 1: Choose a dataset with multi-class classification
# Dataset 2: Choose a dataset with binary classification

print("Loading datasets from OpenML...")
print("=*60

# TODO: Load Dataset 1
dataset1 = fetch_openml(name="adult", as_frame=True)
```

```
# TODO: Load Dataset 2
dataset2 = fetch_openml(name="letter", as_frame=True)

print("Datasets loading section complete.")
```

Loading datasets from OpenML...

=====

Datasets loading section complete.

```
In [32]: def summarize_dataset(data, target, name):
    """Summarize attributes of a dataset."""
    # TODO: Implement dataset summarization

    # 1. Count features
    n_features = data.shape[1]

    # 2. Count instances
    n_instances = data.shape[0]

    # 3. Count classes
    n_classes = len(np.unique(target))

    # 4. Count numerical vs categorical features
    num_features = data.select_dtypes(include=['int64', 'float64']).shape[1]
    cat_features = data.select_dtypes(include=['object', 'category']).shape[1]

    summary = {
        'Dataset': name,
        'Number of Instances': n_instances,
        'Number of Features': n_features,
        'Number of Classes': n_classes,
        'Number of Numerical Features': num_features,
        'Number of Categorical Features': cat_features
    }

    feature_types = {
        'Numerical': num_features,
        'Categorical': cat_features
    }

    return summary, feature_types

# TODO: Call summarize_dataset for both datasets and display results

X1, y1 = dataset1.data, dataset1.target
X2, y2 = dataset2.data, dataset2.target

summary1, types1 = summarize_dataset(X1, y1, "Adult Income")
summary2, types2 = summarize_dataset(X2, y2, "Letter")

print("Dataset Summaries")
```

```
print("=*60)

for summary in [summary1, summary2]:
    print("\n")
    for key, value in summary.items():
        print(f"{key}: {value}")
```

Dataset Summaries

```
Dataset: Adult Income
Number of Instances: 48842
Number of Features: 14
Number of Classes: 2
Number of Numerical Features: 2
Number of Categorical Features: 12
```

```
Dataset: Letter
Number of Instances: 20000
Number of Features: 16
Number of Classes: 26
Number of Numerical Features: 16
Number of Categorical Features: 0
```

Part 2: Training and Evaluation [15 points]

Task:

- Split 80% training / 20% test
- Generate 10 subsets by randomly subsampling 10%, 20%, ..., 100% of training set
- Train Random Forest and Gradient Boosting classifiers
- Measure training time and test accuracy
- Generate learning curves and training time curves

```
In [26]: def prepare_data(data, target):
    """Prepare data for training - handle categorical variables and encoding
    # TODO: Handle encoding of categorical features and target labels

    X = data.copy()

    # One-hot encode categorical columns (if any)
    cat_cols = X.select_dtypes(include=["object", "category"]).columns
    if len(cat_cols) > 0:
        X = pd.get_dummies(X, columns=cat_cols, drop_first=False)
```

```
# Encode target labels
le = LabelEncoder()
y = le.fit_transform(np.array(target))

return X, y

return data, target
```

```
In [27]: def run_experiment(X, y, dataset_name, random_state=42):
    """
    Run the training experiment with 10 different training set sizes.

    Returns:
        results: dict containing accuracies and training times for both
    """
    # TODO: Split data into 80% training and 20% test

    X_train, X_test, y_train, y_test = train_test_split(
        X, y,
        test_size=0.2,
        random_state=random_state,
        stratify=y
    )

    # TODO: Initialize results dictionary to store metrics
    results = {
        "dataset": dataset_name,
        "train_sizes_pct": [],
        "train_sizes_n": [],
        "rf_train_time": [],
        "rf_test_acc": [],
        "gb_train_time": [],
        "gb_test_acc": []
    }

    # Training percentages
    percentages = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 15, 20, 25, 30, 35,

    print(f"\n{'='*70}")
    print(f"Dataset: {dataset_name}")
    print(f"{'='*70}")

    for pct in percentages:
        if pct == 100:
            X_sub, y_sub = X_train, y_train
            subset_size = X_train.shape[0]
        else:
            frac = pct / 100.0
            subset_size = int(np.floor(frac * X_train.shape[0]))
            subset_size = max(1, subset_size)
```

```
X_sub, _, y_sub, _ = train_test_split(  
    X_train, y_train,  
    train_size=subset_size,  
    random_state=random_state,  
    stratify=y_train  
)  
  
# TODO: Train Random Forest and measure time/accuracy  
rf = RandomForestClassifier(  
    n_estimators=200,  
    random_state=random_state,  
    n_jobs=-1  
)  
  
t0 = time.perf_counter()  
rf.fit(X_sub, y_sub)  
rf_time = time.perf_counter() - t0  
  
rf_pred = rf.predict(X_test)  
rf_acc = accuracy_score(y_test, rf_pred)  
  
# TODO: Train Gradient Boosting and measure time/accuracy  
gb = GradientBoostingClassifier(  
    random_state=random_state  
)  
  
t0 = time.perf_counter()  
gb.fit(X_sub, y_sub)  
gb_time = time.perf_counter() - t0  
  
gb_pred = gb.predict(X_test)  
gb_acc = accuracy_score(y_test, gb_pred)  
  
# TODO: Store results  
results["train_sizes_pct"].append(pct)  
results["train_sizes_n"].append(subset_size)  
results["rf_train_time"].append(rf_time)  
results["rf_test_acc"].append(rf_acc)  
results["gb_train_time"].append(gb_time)  
results["gb_test_acc"].append(gb_acc)  
  
print(f"\nTrain subset: {pct}% (n={subset_size})")  
print(f" RF -> time: {rf_time:.4f}s | test acc: {rf_acc:.4f}")  
print(f" GB -> time: {gb_time:.4f}s | test acc: {gb_acc:.4f}")  
  
return results
```

```
In [28]: def plot_results(results, dataset_name):  
    """  
    Generate learning curves and training time curves.  
    """  
    fig, axes = plt.subplots(1, 2, figsize=(14, 5))
```

```

# TODO: Extract data from results dictionary

train_sizes = results["train_sizes_n"]

rf_acc = results["rf_test_acc"]
gb_acc = results["gb_test_acc"]

rf_time = results["rf_train_time"]
gb_time = results["gb_train_time"]

# TODO: Plot Learning Curves (Accuracy vs Data Size) on axes[0]

axes[0].plot(train_sizes, rf_acc, marker='o', label='Random Forest')
axes[0].plot(train_sizes, gb_acc, marker='o', label='Gradient Boosting')

axes[0].set_xlabel("Training Data Size (Number of Samples)")
axes[0].set_ylabel("Test Accuracy")
axes[0].set_title(f"{dataset_name} - Learning Curve")
axes[0].grid(True, alpha=0.3)
axes[0].legend()

# TODO: Plot Training Time Curves (Time vs Data Size) on axes[1]

axes[1].plot(train_sizes, rf_time, marker='o', label='Random Forest')
axes[1].plot(train_sizes, gb_time, marker='o', label='Gradient Boosting')

axes[1].set_xlabel("Training Data Size (Number of Samples)")
axes[1].set_ylabel("Training Time (seconds)")
axes[1].set_title(f"{dataset_name} - Training Time Curve")
axes[1].grid(True, alpha=0.3)
axes[1].legend()

plt.tight_layout()
plt.show()

```

In [21]: # TODO: Run Experiment and Plot for Dataset 1

```
X1, y1 = prepare_data(dataset1.data, dataset1.target)
results1 = run_experiment(X1, y1, "Adult Census Income")
plot_results(results1, "Adult Census Income")
```

```
=====
Dataset: Adult Census Income
=====
```

```
Train subset: 1% (n=390)
RF -> time: 0.3905s | test acc: 0.8215
GB -> time: 0.1973s | test acc: 0.8258
```

```
Train subset: 2% (n=781)
RF -> time: 0.5200s | test acc: 0.8324
GB -> time: 0.2580s | test acc: 0.8403
```

Train subset: 3% (n=1172)
RF -> time: 0.5821s | test acc: 0.8319
GB -> time: 0.3564s | test acc: 0.8446

Train subset: 4% (n=1562)
RF -> time: 0.6540s | test acc: 0.8341
GB -> time: 0.4199s | test acc: 0.8474

Train subset: 5% (n=1953)
RF -> time: 0.7565s | test acc: 0.8381
GB -> time: 0.5115s | test acc: 0.8505

Train subset: 6% (n=2344)
RF -> time: 0.8397s | test acc: 0.8364
GB -> time: 0.6113s | test acc: 0.8515

Train subset: 7% (n=2735)
RF -> time: 0.9202s | test acc: 0.8350
GB -> time: 0.7079s | test acc: 0.8524

Train subset: 8% (n=3125)
RF -> time: 1.4518s | test acc: 0.8370
GB -> time: 1.5460s | test acc: 0.8532

Train subset: 9% (n=3516)
RF -> time: 1.1119s | test acc: 0.8405
GB -> time: 0.8615s | test acc: 0.8523

Train subset: 10% (n=3907)
RF -> time: 1.2385s | test acc: 0.8411
GB -> time: 1.2644s | test acc: 0.8543

Train subset: 15% (n=5860)
RF -> time: 2.2518s | test acc: 0.8395
GB -> time: 1.4011s | test acc: 0.8549

Train subset: 20% (n=7814)
RF -> time: 2.9948s | test acc: 0.8381
GB -> time: 2.1520s | test acc: 0.8575

Train subset: 25% (n=9768)
RF -> time: 2.6010s | test acc: 0.8387
GB -> time: 2.3035s | test acc: 0.8581

Train subset: 30% (n=11721)
RF -> time: 3.1100s | test acc: 0.8374
GB -> time: 4.0757s | test acc: 0.8576

Train subset: 35% (n=13675)
RF -> time: 3.6635s | test acc: 0.8365
GB -> time: 3.6826s | test acc: 0.8580

Train subset: 40% (n=15629)
 RF -> time: 5.7562s | test acc: 0.8379
 GB -> time: 3.6658s | test acc: 0.8589

Train subset: 50% (n=19536)
 RF -> time: 5.8054s | test acc: 0.8383
 GB -> time: 5.3138s | test acc: 0.8567


Train subset: 60% (n=23443)
 RF -> time: 6.8519s | test acc: 0.8388
 GB -> time: 6.4679s | test acc: 0.8578

Train subset: 70% (n=27351)
 RF -> time: 9.4330s | test acc: 0.8371
 GB -> time: 6.4586s | test acc: 0.8575

Train subset: 80% (n=31258)
 RF -> time: 10.6118s | test acc: 0.8364
 GB -> time: 9.2903s | test acc: 0.8576

Train subset: 90% (n=35165)
 RF -> time: 11.7003s | test acc: 0.8353
 GB -> time: 8.3366s | test acc: 0.8583

Train subset: 100% (n=39073)
 RF -> time: 13.8843s | test acc: 0.8348
 GB -> time: 12.2281s | test acc: 0.8568



```
In [34]: # TODO: Run Experiment and Plot for Dataset 2
X2, y2 = prepare_data(dataset2.data, dataset2.target)
results2 = run_experiment(X2, y2, "Letter")
plot_results(results2, "Letter")
```

=====

Dataset: Letter

=====

Train subset: 1% (n=160)
 RF -> time: 1.6822s | test acc: 0.5627
 GB -> time: 7.0862s | test acc: 0.4647

Train subset: 2% (n=320)
RF -> time: 0.5152s | test acc: 0.7013
GB -> time: 3.5189s | test acc: 0.6200

Train subset: 3% (n=480)
RF -> time: 0.4931s | test acc: 0.7508
GB -> time: 5.5821s | test acc: 0.6880

Train subset: 4% (n=640)
RF -> time: 0.5562s | test acc: 0.7728
GB -> time: 4.8102s | test acc: 0.7222

Train subset: 5% (n=800)
RF -> time: 0.6300s | test acc: 0.7910
GB -> time: 6.7519s | test acc: 0.7485

Train subset: 6% (n=960)
RF -> time: 0.6625s | test acc: 0.8055
GB -> time: 5.9367s | test acc: 0.7722

Train subset: 7% (n=1120)
RF -> time: 1.2132s | test acc: 0.8223
GB -> time: 7.2460s | test acc: 0.7817

Train subset: 8% (n=1280)
RF -> time: 0.7357s | test acc: 0.8320
GB -> time: 8.4726s | test acc: 0.8053

Train subset: 9% (n=1440)
RF -> time: 0.7896s | test acc: 0.8455
GB -> time: 8.3851s | test acc: 0.8143

Train subset: 10% (n=1600)
RF -> time: 1.3551s | test acc: 0.8560
GB -> time: 8.2049s | test acc: 0.8227

Train subset: 15% (n=2400)
RF -> time: 1.0130s | test acc: 0.8885
GB -> time: 12.5424s | test acc: 0.8522

Train subset: 20% (n=3200)
RF -> time: 1.8251s | test acc: 0.8882
GB -> time: 16.4928s | test acc: 0.8590

Train subset: 25% (n=4000)
RF -> time: 1.4256s | test acc: 0.9070
GB -> time: 18.1020s | test acc: 0.8738

Train subset: 30% (n=4800)
RF -> time: 1.9816s | test acc: 0.9185
GB -> time: 22.2104s | test acc: 0.8780

Train subset: 35% (n=5600)
 RF -> time: 1.7883s | test acc: 0.9300
 GB -> time: 25.0529s | test acc: 0.8970

Train subset: 40% (n=6400)
 RF -> time: 2.4211s | test acc: 0.9333
 GB -> time: 28.8041s | test acc: 0.8892

Train subset: 50% (n=8000)
 RF -> time: 2.2991s | test acc: 0.9427
 GB -> time: 34.2108s | test acc: 0.9028


Train subset: 60% (n=9600)
 RF -> time: 3.2783s | test acc: 0.9500
 GB -> time: 40.5696s | test acc: 0.9070

Train subset: 70% (n=11200)
 RF -> time: 3.0583s | test acc: 0.9550
 GB -> time: 47.4384s | test acc: 0.9093

Train subset: 80% (n=12800)
 RF -> time: 3.4097s | test acc: 0.9593
 GB -> time: 53.1019s | test acc: 0.9133

Train subset: 90% (n=14400)
 RF -> time: 3.8763s | test acc: 0.9640
 GB -> time: 60.1501s | test acc: 0.9190

Train subset: 100% (n=16000)
 RF -> time: 4.0942s | test acc: 0.9673
 GB -> time: 65.6471s | test acc: 0.9195



Part 3: Analysis and Observations [5 points]

Task: Write three main observations about:

1. Scaling of training time

2. Comparison of accuracy between the two classifiers
3. Learning curve behavior

Your Observations:

Observation 1: Training Time Scaling

- Training time increases approximately linearly with the size of the training data for both classifiers, but the magnitude differs significantly across datasets.
- For the Adult dataset, both Random Forest (RF) and Gradient Boosting (GB) scale moderately with data size. At 100% of training data (~39k samples), RF requires ~13.9 seconds while GB requires ~12.2 seconds. Their training times are comparable, with RF slightly slower at large sizes.
- For the Letter dataset, the difference is much more pronounced. At full training size (~16k samples), RF takes only ~4.1 seconds, whereas GB requires ~65.6 seconds. This indicates that Gradient Boosting scales much more aggressively with data size in multi-class settings (26 classes), while Random Forest remains computationally efficient.
- Overall, Random Forest demonstrates better computational scalability, especially for multi-class problems.

Observation 2: Accuracy Comparison

- On the Adult dataset (binary classification): Gradient Boosting consistently achieves higher test accuracy.
 - GB stabilizes around ~0.857–0.859.
 - RF stabilizes around ~0.835–0.841.
- Thus, Gradient Boosting outperforms Random Forest in terms of predictive accuracy for the binary classification problem.
- On the Letter dataset (26-class classification): Random Forest significantly outperforms Gradient Boosting.
 - RF reaches ~0.967 test accuracy at full data.
 - GB plateaus around ~0.919.
- In multi-class settings with many classes, Random Forest achieves substantially better accuracy.
- Therefore:
 - Binary dataset: Gradient Boosting performs better.

- Multi-class dataset: Random Forest performs better.

Observation 3: Learning Curve Behavior

- For both datasets, accuracy improves rapidly at small training sizes and then plateaus as more data is added.
- On the Adult dataset, accuracy improvements diminish after ~20–30% of training data, indicating the model reaches its performance limit relatively early. The learning curves flatten, suggesting low additional benefit from further increasing data size.
- On the Letter dataset, accuracy steadily increases with more data, especially for Random Forest. The improvement remains visible even beyond 50% of training data, indicating that additional data continues to reduce variance and improve generalization.
- This demonstrates: Adult dataset (Binary Classification) shows early saturation. Letter dataset (Multi-Class Classification) benefits more consistently from additional training data. Random Forest appears more data-efficient in high-class-count settings.